

Introduction to Neural Networks

June 22, 2020

Neural Networks try to resemble the way humans make decisions and the way neurons in the human brain take in data and act on input.

Introduction

A neuron usually has a couple parts to it. There are the **dendrites**, which are in charge of receiving incoming signals from other neurons. The signals then travel through the cell body and the **axons** all the way to the connections at the end of the axons. These neurons are all connected to many other neurons each, so any attempt to model these connections will necessarily use up a lot of connections.

In computer science, we have something resembling a neuron, except that the way we model decisions is a bit different.

Neural networks are usually represented by a connected graph of neurons, usually organized in layers. There are a couple types:

1. **Feedforward**: In these types, the signals only flow in one direction.
2. **Recurrent**: In recurrent neural networks, there is feed forward functionality, but there can also be feedback to the neurons.

In a neural network, we would usually have k output layers, which might represent the confidence with which we classify or predict a certain value for our model. Let the final value of our model be $y(x)$ for some input vector x . The error value for a neural network would then be

$$e = d - y$$

which we of course want to minimize.

Perceptron Algorithm

The perceptron algorithm is an algorithm which can learn the classification from the data. How does it do this? It only updates the set of weights when it makes a mistake in classifying or regressing(?) the testing data.

Originally, with a single neuron, the algorithm could only handle the cases that were **linearly separable**. For example, training a model to learn the behavior of the XOR function is not possible with only a single neuron, again since the training data is not linearly separable. However, with a simple two-layer network we can easily model the behavior of XOR.

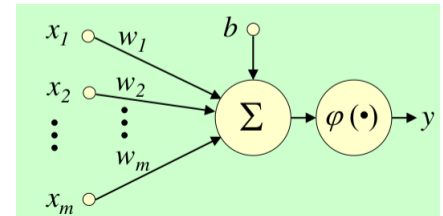


Figure 1: We take a set of inputs $x_1 \dots x_n$ and pass each of these through a corresponding **weight function** $w_1 \dots w_n$. In the actual node body, we take the sum of the weight functions connecting the current neuron to the ones in the previous layer. This sum Σ is then passed through the final **activation function** ϕ , which gives the final output for this neuron.

So, for a single neuron, updating the values of \mathbf{w} is quite easy, since we only have a single value to update once we have estimated the model. How do we do this when we have multiple values in \mathbf{w} ? We have to use **Back-Propagation**.

Back-Propagation tries to maximize the values starting at the output nodes. The error value then moves back through the layers, and we calculate the error for the individual neurons in a backward fashion. An iteration in which all the testing samples are seen at least once is called an **epoch**.

Once we do the back-propagation, how do we know when to stop? We can have several criteria:

1. Stop when the cost function is lower than a threshold (when it's "good enough")
2. Also can stop when the gradient is lower than a certain function.
3. When the **learning curve** over several iterations remains stable.

Remember the **activation function** tells us how confident we are in our prediction, and it lies between 0 and 1, for not confident and confident, respectively. What functions can map the entire number interval into 0 and 1? The old standard was the **sigmoidal function**, which is

$$f(v) = \frac{1}{1 + e^{-av}}$$

There appears to be some issues with this function. Although it is **continuously differentiable**, the gradients apparently go to infinity pretty quickly? This function used to be the standard, but not so much anymore.

Now, the standard for activation functions has become **REctified Linear Units (RELU)**. This function takes the form

$$f(v) = \tanh(av)$$

Both of these functions are **non-linear**.

Another question remains: What should our learning rate be? We could either choose to make it larger or smaller, and there are consequences for both. If we make it smaller, then the algorithm will take slower, since to find the global maximum it will have to be adjusted several times. However, the learning rate φ will proceed along a smoother trajectory to the final learning rate.

For a larger φ , the effects are the opposite: there is a faster training cycle but more oscillation from the different epochs.

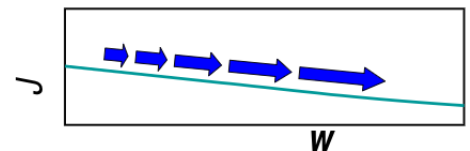


Figure 2: