Deep Learning and CNNs

Andrés Ponce

June 5, 2020

Deep Learning methods allow us to perform other functions such as object recognition obtained from image datasets.

Conventional Approach to Object Recognition

The key steps in object recognition involve **feature extraction** from several images. However, features may look different when recognizing different types of objects.

We could have a mdoel with **hand-crafted** feature detection. Here, we could define by hand the kind of features we would like to see in our model. After we go through the feature extractor we could have a simple classifier.

Another approach is instead to have a **trainable feature extractor**, along with a trainable classifier.

Neural Networks

If we recall, **neural networks** are a series of layers, composed of individual **neurons** and which connect and communicate to neurons the layer before or after it.

Training the neural network consists in trying to minimzize the loss function. For classification, the loss function might be the distance between the correct classification and our predicted values. The formula might look something like:

$$L(\theta) = \sum_{i=1}^{N} \|\mathbf{y}_i - g_w(\mathbf{x_i})\|^2$$

To minimize it, we use the technique called **gradient descent**. This technique involves finding the values for which the **derivative** of the loss function decreases. Even within gradient descent, there are some different types we might wish to use. For example, **steepest descent** uses batch processing to change the parameter vector $\mathbf{w}$. [1]

Another method for gradient descent involves selecting a subset of the dataset to use when measuring error. This is called **stochastic** gradient descent. Here we use a formula such as

$$E(w) = \frac{1}{2} \sum_{n=1}^{N} \{y(x_n, w) - t_n\}^2$$

also known as **sum of squares**. [2]

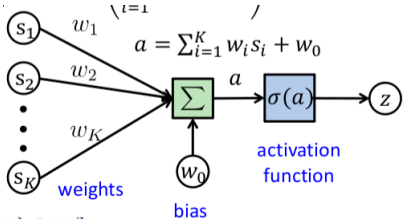

Figure 1: In neural networks, we have a set $\mathbf{s}$ of input values, which are then summed up through the different layers. Afterwards, we pass it through an **activation function** $\sigma(a)$, which then leads to the class we predict it into.

[1] Remember batch methods operate on the whole training set to modify weights.

[2] Sum of squares errors are more useful for problems resembling regression.