

Ensemble Methods

Andrés Ponce

May 7, 2020

Ensemble Methods refer to when we use an aggregate of L different models to find more accurate predictions for our data. There are several ways of implementing a selection from several methods: adaboost, decision trees, bagging, random forests, etc...

Boosting

The straightforward approach to using multiple models is to take the average value of each of the models that we use. Supposing we want to use M models, and we want to predict the value given by a certain function, we could say

$$y_{COM} = \frac{1}{M} \sum_{m=1}^M y_m(x)$$

where *COM* refers to the *decision by committee* given by the equations. We take the average value of the M models to form our final result. When calculating the error, we also take the aggregate error of all the models.

Boosting then refers to the building of a strong classifier for a problem by using several weaker classifiers. When boosting, we again can have M models whose scores are summed up.

The type of boosting called **Adaptive Boosting** will pass the input through the models sequentially, and *train* the models sequentially. Then, when training the next model, items that were misclassified will be given greater weight. This may result in an overall model which performs very well even if the models themselves perform as well as random ones.

The process for AdaBoost looks something like:

1. Given inputs $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X, y_i \in \{-1, +1\}$, initialize their distributions $D_t(i) = \frac{1}{m}, i = 1, \dots, m$ ¹
2. Find the classifier h_t which minimizes the error w.r.t. D_t .²
3. Calculate the weight classifier $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$
4. Update our distribution³

$$D_{t+1}(i) = \frac{D_t \exp[-\alpha_t y_i h_t(x_i)]}{Z_t}$$

Once we have the final classifier $H(x)$, comprised of the sum of the models with the adjusted weights and distributions, we can calculate




Figure 1: Each of the models contributes to how we pass the input set to the next model to be trained. At the end, we also use each model's results to calculate the final prediction.

¹ The distributions $D_t(i)$ mean the probability distribution of x_i 's at time t

² We can basically take the min of

$$\epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j(x_i)]$$

³ Z_t is basically a value to normalize the distribution.

the **margin** of the classifier on a sample (x, y) . The margin is given by⁴

$$yH(x)$$

The AdaBoost algorithm will try to maximize the margins of our variables by minimizing the **exponential loss** of the predictions.

$$loss_{exp}[H(x)] = E_{x,y}[e^{-yH(x)}]$$

Face Detection using AdaBoost

To perform facial detection on an image, we can use **sliding windows** which apply some model to different portions of the image. However, this approach might be a bit challenging because of the large amount of areas that we would need to detect in a single image. Since the amount of pixels might be very large, we would need to beware of the amount of false-positives that are possible. For this problem, the **Viola-Jones** method is quite useful.

Viola-Jones Face Detector

The Viola-Jones method for detecting faces relies on finding **rectangle features** in images. Since a face can have eyes, nose, etc... we can measure the difference in pixel areas between similar pattern across faces. For example, if we have an idea of how different the eyes and nose should look like, we can use a model similar to Figure 2 to measure that area. If the end result is something which we might expect, then we can feel confident in saying there is a face or other feature inside.

A "feature" in Viola-Jones corresponds to a difference in pixel values. That's why having different models with clear/black parts is useful: it can help us detect different parts in a face. Thus, for a feature value $f(x)$ given an input image vector x can be written as:

$$\text{Feature Value} = \sum (\text{Pixels in white area}) - \sum (\text{Pixels in black area})$$

The next distinguishing feature of Viola-Jones facial recognition involves the **integral image**. This aspect involves keeping a record of the sum of all preceding parts of an image. For an integral image ii , the value $ii(x, y)$ is the sum of all values *above and to the left of* x, y .

This way of storing cumulative image values makes finding the overall value in a rectangular area quite easy. In the end, we only need a constant number of additions to calculate the feature value of any rectangular area in our image.

⁴ The margin tells us when we classified the input correctly, and our confidence with our classification.




Figure 2: In Viola-Jones, there are four types of features of rectangular features. Different models might be applicable to different parts of the face. We have to take the difference of the clear and the black areas. C might be used to measure eyes and nose area; B for eyes and mouth, etc...




Figure 3: The value at $ii(x, y)$ is the sum of all the values that came before it.

Boosting for face detection

So far, we've only used Viola-Jones to detect faces through the integral image. Then how do we use boosting with this approach? Similar to the definition covered before, we can use each one of the different classifiers for the rectangle areas.

$$h(x) = \begin{cases} 1 & p_t f_t > p_t \theta_t \\ 0 & \text{otherwise} \end{cases}$$